


CrossMEP: A Tiered Synthetic Dataset of Multi-Trade MEP Cross-Sections

A support-design problem distribution for learning-based structural support assembly synthesis

Lavinia Pedrollo¹ • David Gvadzabia² • Torben Graeber³ • Martin Fischer¹

¹ Stanford University, CIFE ² Lafayette College ³ Hilti AG

43rd International CIB W78 Conference • IT in Construction • New Delhi, 6–8 October 2026




a generated C8 context: eight services on three rows

A hospital needs about 10,000 support assemblies, each designed by hand



1 unit = 1 structural support assembly (SSA)



$\approx 10,000$

support assemblies in one 200,000 sq ft hospital

20 min – 2 h

of engineering for each assembly, designed one by one


$\approx \frac{1}{4}$

of MEP design effort; on the order of \$600K per project

Approximate practitioner estimates.


A designer turns one cross-section into verified support assemblies

In: one cross-section = a context



Structure: ceiling, floor or wall, with anchor zones
Services: trade, position, size and weight per metre

Out: verified support designs = assemblies




Up to ten designs, ranked by installed cost
Each checked: statics, anchors, connectors, buildability

Context, the brief: the section at one hanger, its structure and services.
This is what CrossMEP contains.

Assembly, the answer: catalog parts that carry the services to the structure. Not in CrossMEP: a method produces it.

The synthesis problem as defined at ISARC 2026.

Synthesizing the assembly from a catalog and a context has resisted automation



(a) Full MEP route



(b) One CrossMEP context

Design starts after coordination: at each hanger the designer works from the section across the run, not from the whole model.

Where practice stands

Tools coordinate the model and check a design. Choosing the topology and the parts, from a fixed catalog, is still done by hand, one section at a time.

1. Geometric judgment

Engineers choose an assembly's topology from experience.


2. Rule explosion

Code, load and material constraints interact across thousands of combinations.

3. Catalog volatility



Catalogs change faster than rule-based systems can be rewritten.

Whatever the method, it needs problems to learn from and a fixed set to compare on



Whatever sits in the black box, it needs **problems to learn from**, a **fixed set to compare on** and **problems that look like practice**.


Real designs cover one corner of the design space; CrossMEP covers all of it



CrossMEP at a glance

7,000 contexts in four splits on disjoint seeds	31,484 elements: pipes, cable trays, ducts and conduits; six trades	C1-C8 eight difficulty tiers by exact element count
0 labels the brief, not the answer: no support-assembly information	Traced every constant with its source, or declared a design choice	Open data CC BY 4.0, code MIT; a seeded, deterministic generator


We generate what a designer receives: the section at one hanger



Per element: kind, service and trade · bare size · insulation per side · load per metre × the span it was sized at = load at the support · position along and out from the surface. Per context: slab or wall, substrate, thickness. Absent by design: channel, rods, clamps, anchors — and any “correct” answer, since a feasible support depends on the catalog you build from.

Sections follow trade practice and measured spacing, not random shapes

- 1 Tier.** Cn fixes the exact element count, $n = 1 \dots 8$.
- 2 Surface.** Ceiling or wall (78 / 22 %); concrete slab or wall, 150–300 mm thick.
- 3 Services.** Grouped as trades run: hot + cold pairs, flow + return, chilled banks, sprinkler mains, 1–2 trays, ducts, conduit groups of 2–6.
- 4 Rows.** Bulky first: ducts nearest the surface, then trays and conduits, then pipes.
- 5 Spacing.** Gaps drawn from gaps measured on a built project (at least 25 mm); stagger within rows; electrical above wet on walls.
- 6 Check.** Every pair at least 25 mm clear; rows centred; sizes, insulation, loads and positions recorded.



One generated benchmark context, three rows by priority; the gap marked is one draw from the measured distribution. Every context is seeded: the same seed gives the same file, byte for byte.

Sizes, spans and loads come from standards; the choices are declared

Element	Sizes	Load basis
Pipes DN15–100	EN 10220 / EN 10255 medium series	steel + water × ASME B31.1 water-service span (published points only)
Insulation	GEG Anlage 8 (heated); 30/50 mm condensation control (chilled)	—
Cable trays 150–600	IEC 61537 systems, NEMA VE 1 widths	full tray: 50 kg/m at 300 mm datum, × 2.0 m (NEMA VE 1 class span)
Ducts	EN 1505 rectangular; EN 1506 round	manufacturer duct-weight table incl. flanges, × 2.4 m (SMACNA)
Conduits Ø20–50	IEC 61386-1, parallel groups of 2–6	steel tube + 40 %-of-bore cable fill, × 2.0 m (IET / BS 7671)

Layout conventions



- Bulky services nearest the slab: ducts, then containment, then pipes
- Trades run in banks: hot + cold pairs, flow + return, conduit groups
- Electrical kept above wet services on walls (drip; BS 7671 528.3.2)
- Clear gaps drawn from a distribution measured on a built project; 25 mm minimum
- Within-row stagger calibrated to measured elevation spread

Three kinds of ground: standards (what each element is and weighs), measured open buildings (how close neighbours sit) and declared design choices (trade mix, surface mix, number of rows). Every constant carries its source and status in the repository, and the log says why each standard was chosen over its alternatives.


7,000 contexts in four splits; the test seeds are never trained on

split	contexts	elements	seed
train	5,000	22,500	1000
validation	500	2,242	2000
test	500	2,242	3000
benchmark	1,000	4,500	42
total	7,000	31,484	

- Tiers assigned round-robin: 625 per tier in train, 125 in the benchmark
- Four in five contexts hang from a ceiling, one in five from a wall
- Six trades: electrical 10,494, domestic water 5,725, heating 5,788, chilled water 3,987, sprinkler 2,073, ventilation 3,417 elements
- Plain JSON, with a JSON Schema, Croissant metadata and a datasheet




Difficulty is one number: tier Cn holds exactly n elements




One benchmark context per tier; dashed rings are insulation. Within a tier, kinds, trades, services, surfaces and stacking all vary, so element count is the one controlled difficulty axis. C1 is the most common support in any building; C8 is a congested rack.

Higher tiers are heavier and tighter, as designed





A check that the dataset behaves as designed: the tier fixes the element count; load and congestion follow from the rules. Boxes: benchmark, 125 contexts per tier (IQR, whiskers 1.5 IQR); diamonds: medians of 2,000 generated contexts per tier. Load rises with the tier (rank correlation +0.46; the population medians rise at every step) and the closest gap narrows at every step (-0.38; 149 → 55 mm in the population).

Spacing holds up on a clinic the generator never saw



The gap distribution was fitted on the duplex; the clinic was held out. Clear gaps between side-by-side pipes, on sections every 250 mm. Clinic (797 pipe pairs): 32 mm from the generated gaps (95 % CI 17–48), more than the 8 mm a perfect generator would show, and about as close as the duplex’s own two models are to each other (26 mm). Models: buildingSMART Duplex Apartment and Medical-Dental Clinic, CC BY 4.0.

A two-size catalog attaches 1 pipe in 7; the dataset shows what to cover



14.5 %

of benchmark pipes attachable with the paper's two sizes (95 % CI 12.7–16.2); every one is DN40

38.1 %

with two sizes placed where the pipes are; 76 % with six, and every pipe with 15

0 of 2,012

trays, ducts and conduits: a pipe-clamp catalog defines no attachment. Load never binds (1.97 vs 2.5 kN)

Set the parameters and the generator returns the designer's brief

SERVICES AT THE HANGER

TIER = NUMBER OF SERVICES



SEED

By tier Exact mix

C1 C2 C3 C4 C5 C6 C7 C8

< 0 >

Generate another



Scan to open the studio


In: the parameters (tier C7, seed 0). Out: section A-A at the support, as the support designer receives it (scale 1:20), and the same run in 3D. Every section is a stored output of the released generator, with the Python call that reproduces it.

Train on the generator, evaluate on the benchmark, report per tier

```
import crossmep.tasks as cm
data = cm.load("benchmark") # 1,000 contexts
cm.filter_contexts(data, pipes=2, trays=1)

from crossmep.evaluate import score, compare
score(my_results, data) # per tier, 95 % CI
compare(mine, baseline, data) # paired test
```

Loading, metrics and scoring need only the Python standard library; the generator needs NumPy. Command line: `python -m crossmep generate | validate | evaluate`.



Interactive gallery: filter the benchmark by tier, kind, trade and surface.

Train

on the generator: unlimited seeds, a curriculum from C1 up to C8

Evaluate

on the benchmark: 125 held-out contexts per tier, the same for every method

Report

per tier with 95 % intervals, paired tests between methods

One section at one support today; a checker and real data come next

Scope

- One cross-section at one support: routing, branches and elevation changes are outside it; so is support spacing — the span is recorded per element, so a method can vary it
- Realism checked on sections of two open buildings, a residential duplex and a medical clinic; congested racks rest on coordination practice and standards
- Trade mix and tier composition are design choices, not survey results
- Synthetic: not for the structural design of real installations

Next

- A public checker for support designs, with best-known costs from an exact solver on the small tiers
- The catalog as an input: generated catalogs, and tests on catalogs a method has never seen
- A real test set: sections from commercial projects with supports designed by engineers

If you coordinate services or design supports, open the gallery, filter for the scenes you know, and tell us what looks wrong. Every rule is one documented constant away from being changed.



Support design now has open problems to learn from


7,000 support-design problems · every constant sourced or declared
pipe gaps checked on two open buildings · open data and code

Data CC BY 4.0, code MIT. Not for the structural design of real installations.




github.com/lavinia-ped/CrossMEP_Dataset
lavinia@stanford.edu

A wall with six pipes, and a slab with a duct, a tray and conduits



Tier C6, seed 0, section 5: six pipes fixed to a wall; offsets from the wall face, the closest clear gap 25 mm.



Exact mix: 2 pipes, 1 tray, 1 duct, 2 conduits on a slab; levels below the soffit and the load at the support.

Median load rises from 0.30 kN at C1 to 2.04 kN at C8 as the gap narrows

tier	contexts	median clear gap (mm)	median load (kN)	median width (mm)
C1	125	—	0.30	140
C2	125	130	0.24	431
C3	125	89	1.04	776
C4	125	77	0.94	1037
C5	125	72	1.22	1275
C6	125	69	1.36	1494
C7	125	61	2.24	1861
C8	125	51	2.04	2225

4,500 elements

2,488 pipes · 968 conduits · 559 trays · 485 ducts

32 mm

Wasserstein-1, generated to measured pipe gaps of the clinic
(95 % CI 17–48)

–0.38 / +0.46

rank correlation of tier with the closest gap / with the load

Against the clinic the generator is 32 mm off; the small duplex is 76 to 88

pipe gaps	pairs	sections	median (mm)	< 25 mm
generated (benchmark)	1,690	—	216	0 %
Clinic, Plumbing	797	3,833	203	2 %
Duplex, MEP	74	204	177	23 %
Duplex, Plumbing	41	112	114	32 %

- Sections every 250 mm; runs within 400 mm of height form a row; gaps between bare pipe surfaces, below 600 mm
- Robust: across 11 settings of the cut (spacing, row band, axis tolerance, run length) generated ↔ clinic stays at 31–41 mm
- Noise floor: the distance a perfect generator would still show at that sample size
- The duplex samples are small, hence the wide intervals

Wasserstein-1 (mm)	by length	each pair once	noise floor
generated ↔ Clinic	32 (17–48)	18 (13–27)	8
generated ↔ Duplex MEP	76 (61–103)	97 (78–120)	23
generated ↔ Duplex Plumbing	88 (62–132)	104 (89–127)	32
Clinic ↔ Duplex MEP	71 (54–101)	88 (70–118)	—
Clinic ↔ Duplex Plumbing	85 (55–137)	94 (77–121)	—
Duplex MEP ↔ Duplex Plumbing	26 (24–124)	25 (21–108)	—

Intervals: 95 % bootstrap over contexts (generated) and pipe pairs (measured). Models: buildingSMART Duplex Apartment and Medical-Dental Clinic, CC BY 4.0; derived measurements and attribution in [verify/measured/](#).

The diameter rule moves catalog coverage more than tier or split does

tier	pipes	benchmark % (95 % CI)	population % (95 % CI)
C1	66	16.7 (8.3–26.6)	12.7 (10.7–14.7)
C2	132	11.4 (5.0–18.4)	13.5 (11.8–15.4)
C3	212	15.6 (10.5–21.1)	13.8 (12.4–15.2)
C4	284	16.9 (11.8–22.1)	13.2 (12.0–14.4)
C5	365	11.5 (7.3–16.0)	14.7 (13.5–15.9)
C6	366	10.7 (7.2–14.5)	14.3 (13.3–15.4)
C7	517	17.2 (13.4–21.1)	14.9 (13.9–16.0)
C8	546	15.2 (11.7–19.1)	14.4 (13.5–15.4)

attach diameter rule	bins ± 0 mm	bins ± 0.5 mm	bins ± 1 mm
service (paper)	14.5	19.8	19.8
bare	14.5	22.0	22.0
insulated	7.2	8.2	8.2

% of benchmark pipes attachable when both bin edges are widened by the stated amount. DN100 (114.3 mm) lies 0.3 mm above the 108–114 mm bin, so the headline moves with the diameter rule and the edge tolerance.

split	pipes	% attachable (95 % CI)
train	12,590	14.4 (13.7–15.2)
val	1,295	13.2 (11.2–15.6)
test	1,200	13.8 (11.5–16.4)
benchmark	2,488	14.5 (12.7–16.2)

- Intervals: cluster bootstrap over contexts (2,000 resamples), because the elements of one context are correlated
- Population: 2,000 freshly generated contexts per tier; overall 14.3 % (95 % CI 13.8–14.7)
- Only size and capacity are tested, so coverage is an upper bound on what can be installed
- Best possible share with k sizes: exact optimum over the distinct attach diameters

A context is stored as plain JSON: elements, sizes, loads, positions

```
{
  "context_id": "mep_9", "tier": "C2",
  "surface": {
    "kind": "ceiling", "substrate": "concrete_C25", "thickness_mm": 200.0
  },
  "n_elements": 2, "n_levels": 1, "total_load_kN": 0.76, "bundle_width_mm": 1527.0,
  "elements": [
    {
      "kind": "duct", "service": "duct", "trade": "ventilation", "shape": "rect",
      "label": "1000x500 duct", "width_mm": 1000.0, "height_mm": 500.0, "insulation_mm": 0.0,
      "load_kN": 0.72, "span_m": 2.4, "load_kN_per_m": 0.3002,
      "level": 0, "along_mm": -263.5, "out_mm": 340.0
    },
    {
      "kind": "pipe", "service": "domestic_cold", "trade": "domestic", "shape": "round",
      "label": "DN20", "width_mm": 26.9, "height_mm": 26.9, "insulation_mm": 0.0,
      "load_kN": 0.04, "span_m": 2.1, "load_kN_per_m": 0.0189,
      "level": 0, "along_mm": 750.0, "out_mm": 341.3
    }
  ]
}
```

- Units: millimetres, kilonewtons, metres (span)
- Widths are bare; insulation_mm is per side
- $\text{load_kN} = \text{load_kN_per_m} \times \text{span_m}$
- along_mm, out_mm: position along and out from the surface
- level: the generative row (metadata)